

[16] Survey of Hallucination in Natural Language Generation

- **Authors:** Z. Ji, N. Lee, R. Frieske, T. Yu, D. Su, Y. Xu, E. Ishii, Y. J. Yeung, A. Del Luceno, & P. Fung
 - **Venue & Year:** ACM Computing Surveys (2023), 55(12), 1–38
 - **Category:** Hallucination & Grounding
 - **Associated Full Paper PDF:** [ji_2023_hallucination_survey.pdf](#)
-

1-Minute Executive Takeaway

Key Finding: Comprehensive taxonomy defining intrinsic and extrinsic hallucination and establishing retrieval as the primary countermeasure.

Core Research Objective

To provide a comprehensive taxonomy, detection methodology, and mitigation survey of hallucinations in natural language generation.

Methodology & Architecture

Categorized hallucinations into intrinsic (conflicting with source evidence) and extrinsic (unverifiable from source text), surveying retrieval-based, decoding-based, and post-processing mitigation techniques.

Key Empirical Findings & Evidence

Concluded that conditioning generation on retrieved evidence chunks and enforcing structured schema constraints are among the most effective mechanisms to curb factual fabrication.

Direct Relevance to Our Kwiz Paper

Cited in Section 4.2 for our Context Groundedness (CG) evaluation rubric. Used to formally define factual hallucination against course slide chunks and justify our two-tier validation engine.