

[15] Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena

- **Authors:** L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, & I. Stoica
 - **Venue & Year:** Advances in Neural Information Processing Systems (NeurIPS 2023), 36, 46595–46623
 - **Category:** Automated Evaluation Protocol
 - **Associated Full Paper PDF:** [zheng_2023_judging_llm_as_a_judge.pdf](#)
-

1-Minute Executive Takeaway

Key Finding: The benchmark study validating LLM-as-a-Judge as a scientifically rigorous evaluation methodology matching human agreement.

Core Research Objective

To examine whether frontier large language models (such as GPT-4) can act as automated, scalable, and reliable evaluators for natural language generation benchmarks.

Methodology & Architecture

Benchmarked LLM judge agreement against human expert consensus across pairwise and Likert evaluations on MT-Bench and Chatbot Arena, analyzing position bias, verbosity bias, and self-enhancement bias.

Key Empirical Findings & Evidence

Demonstrated that strong frontier LLM judges achieve >80% agreement with human experts, matching human-to-human inter-rater reliability levels.

Direct Relevance to Our Kwiz Paper

Methodological foundation for Section 4.2. We adopted Zheng et al.'s LLM-as-a-Judge framework (deploying GPT-4o and Gemini 2.5 Flash as multi-evaluators) alongside calibrated human faculty grading to evaluate our 100-item MCQ corpus across five dimensions.